

# **Chapter 1**

## **Introduction**

---

**Distributed Systems**

# Distributed versus Decentralized

What many people state



Centralized



Decentralized



Distributed

# Distributed versus Decentralized

What many people state



Centralized



Decentralized



Distributed

When does a decentralized system become distributed?

# Distributed versus Decentralized

What many people state



Centralized



Decentralized



Distributed

When does a decentralized system become distributed?

- Adding 1 link between two nodes in a decentralized system?

# Distributed versus Decentralized

What many people state



Centralized



Decentralized



Distributed

When does a decentralized system become distributed?

- Adding 1 link between two nodes in a decentralized system?
- Adding 2 links between two other nodes?

# Distributed versus Decentralized

What many people state



Centralized



Decentralized



Distributed

When does a decentralized system become distributed?

- Adding 1 link between two nodes in a decentralized system?
- Adding 2 links between two other nodes?
- In general: adding  $k > 0$  links....?

# Alternative approach

## Two views on realizing distributed systems

- **Integrative view:** connecting existing networked computer systems into a larger a system.
- **Expansive view:** an existing networked computer systems is extended with additional computers

# Alternative approach

## Two views on realizing distributed systems

- **Integrative view**: connecting existing networked computer systems into a larger a system.
- **Expansive view**: an existing networked computer systems is extended with additional computers

## Two definitions

- A **decentralized system** is a networked computer system in which processes and resources are **necessarily** spread across multiple computers.
- A **distributed system** is a networked computer system in which processes and resources are **sufficiently** spread across multiple computers.



## What are Distributed Systems?

A group of separate computers working together to achieve a certain shared goal.

These computers, also known as nodes, collaborate, exchange information via a network, and plan their actions to accomplish a shared objective by sharing resources, tasks, and data.

## Some common misconceptions

### Centralized solutions do not scale

Make distinction between **logically** and **physically** centralized. The root of the Domain Name System:

- logically centralized
- physically (massively) distributed
- decentralized across several organizations

### Centralized solutions have a single point of failure

Generally not true (e.g., the root of DNS). A single point of failure is often:

- easier to manage
- easier to make more robust

# Perspectives on distributed systems

## Distributed systems are complex: take perspectives

- **Architecture**: common organizations
- **Process**: what kind of processes, and their relationships
- **Communication**: facilities for exchanging data
- **Coordination**: application-independent algorithms
- **Naming**: how do you identify resources?
- **Consistency** and **replication**: performance requires of data, which need to be **the same**
- **Fault tolerance**: keep running in the presence of partial failures
- **Security**: ensure authorized access to resources

# What do we want to achieve?

## Overall design goals

- Support sharing of resources
- Distribution transparency
- Openness
- Scalability

# Sharing resources

## Canonical examples


- Cloud-based shared storage and files
- Peer-to-peer assisted multimedia streaming
- Shared mail services (think of outsourced mail systems)
- Shared Web hosting (think of content distribution networks)

## Observation

*"The network is the computer"*

(quote from John Gage, then at Sun Microsystems)


# Distribution transparency



## What is transparency?

*The phenomenon by which a distributed system attempts to **hide** the fact that its processes and resources are **physically distributed across multiple computers**, possibly **separated by large distances**.*

# Distribution transparency



## What is transparency?

*The phenomenon by which a distributed system attempts to **hide** the fact that its processes and resources are **physically distributed across multiple computers**, possibly **separated by large distances**.*

## Observation

Distribution transparency is handled through many different techniques in a layer between applications and operating systems: a **middleware layer**

# Distribution transparency

## Types

Transparency	Description
Access	Hide differences in data representation and how an object is accessed
Location	Hide where an object is located
Relocation	Hide that an object may be moved to another location while in use
Migration	Hide that an object may move to another location
Replication	Hide that an object is replicated
Concurrency	Hide that an object may be shared by several independent users
Failure	Hide the failure and recovery of an object



# Degree of transparency

Aiming at full distribution transparency may be too much

# Degree of transparency

Aiming at full distribution transparency may be too much

- There are communication latencies that cannot be hidden

# Degree of transparency

Aiming at full distribution transparency may be too much

- There are communication latencies that cannot be hidden
- **Completely hiding failures** of networks and nodes is (theoretically and practically) **impossible**
  - You cannot distinguish a slow computer from a failing one
  - You can never be sure that a server actually performed an operation before a crash

# Degree of transparency

## Aiming at full distribution transparency may be too much

- There are communication latencies that cannot be hidden
- **Completely hiding failures** of networks and nodes is (theoretically and practically) **impossible**
  - You cannot distinguish a slow computer from a failing one
  - You can never be sure that a server actually performed an operation before a crash
- Full transparency will **cost performance**, exposing distribution of the system
  - Keeping replicas **exactly** up-to-date with the master **takes time**
  - Immediately flushing write operations to disk for fault tolerance

# Degree of transparency

## Exposing distribution may be good

- Making use of location-based services (finding your nearby friends)
- When dealing with users in different time zones
- When it makes it easier for a user to understand what's going on (when e.g., a server does not respond for a long time, report it as failing).

# Degree of transparency

## Exposing distribution may be good

- Making use of location-based services (finding your nearby friends)
- When dealing with users in different time zones
- When it makes it easier for a user to understand what's going on (when e.g., a server does not respond for a long time, report it as failing).

## Conclusion

Distribution transparency is a nice goal, but achieving it is a different story, and it should often not even be aimed at.

# Openness of distributed systems

## Open distributed system

*A system that **offers components** that can easily be used by, or **integrated into other systems**. An open distributed system itself will often consist of components that originate from elsewhere.*

## What are we talking about?

Be able to interact with services from other open systems, irrespective of the underlying environment:

- Systems should conform to well-defined **interfaces**
- Systems should easily **interoperate**
- Systems should support **portability** of applications
- Systems should be easily **extensible**

# Policies versus mechanisms

## Implementing openness: policies

- What level of consistency do we require for client-cached data?
- Which operations do we allow downloaded code to perform?
- Which QoS requirements do we adjust in the face of varying bandwidth?
- What level of secrecy do we require for communication?

## Implementing openness: mechanisms

- Allow (dynamic) setting of caching policies
- Support different levels of trust for mobile code
- Provide adjustable QoS parameters per data stream
- Offer different encryption algorithms



# On strict separation

## Observation

The stricter the separation between policy and mechanism, the more we need to ensure proper mechanisms, potentially leading to many configuration parameters and complex management.

## Finding a balance

Hard-coding policies often simplifies management, **and** reduces complexity at the price of less flexibility. There is no obvious solution.

# Dependability

## Basics

A **component** provides **services** to **clients**. To provide services, the component may require the services from other components  $\Rightarrow$  a component may **depend** on some other component.

## Specifically

A component  $C$  depends on  $C^*$  if the **correctness** of  $C$ 's behavior depends on the correctness of  $C^*$ 's behavior. (Components are processes or channels.)

# Dependability

## Requirements related to dependability

Requirement	Description
Availability	Readiness for usage
Reliability	Continuity of service delivery
Safety	Very low probability of catastrophes
Maintainability	How easy can a failed system be repaired

# Reliability versus availability

## Reliability $R(t)$ of component $C$

Conditional probability that  $C$  has been functioning correctly during  $[0, t)$  given  $C$  was functioning correctly at the time  $T = 0$ .

## Traditional metrics

- **Mean Time To Failure** ( $MTTF$ ): The average time until a component fails.
- **Mean Time To Repair** ( $MTTR$ ): The average time needed to repair a component.
- **Mean Time Between Failures** ( $MTBF$ ): Simply  $MTTF + MTTR$ .

# Terminology

## Failure, error, fault

Term	Description	Example
Failure	A component is not living up to its specifications	Crashed program
Error	Part of a component that can lead to a failure	Programming bug
Fault	Cause of an error	Sloppy programmer

## On security

### Observation

A distributed system that is not secure, is not dependable

# On security

## Observation

A distributed system that is not secure, is not dependable

## What we need

- **Confidentiality**: information is disclosed only to authorized parties
- **Integrity**: Ensure that alterations to assets of a system can be made only in an authorized way

# On security

## Observation

A distributed system that is not secure, is not dependable

## What we need

- **Confidentiality**: information is disclosed only to authorized parties
- **Integrity**: Ensure that alterations to assets of a system can be made only in an authorized way

## Authorization, Authentication, Trust

- **Authentication**: verifying the correctness of a claimed identity
- **Authorization**: does an identified entity has proper access rights?
- **Trust**: one entity can be assured that another will perform particular actions according to a specific expectation



# Security mechanisms

## Keeping it simple

It's all about **encrypting** and **decrypting** data using **security keys**.

## Notation

$K(data)$  denotes that we **use key  $K$**  to **encrypt/decrypt  $data$** .

# Security mechanisms

## Symmetric cryptosystem

With encryption key  $E_K(data)$  and decryption key  $D_K(data)$ :

if  $data = D_K(E_K(data))$  then  $D_K = E_K$ . Note: encryption and decryption key are the same and should be kept **secret**.

## Asymmetric cryptosystem

Distinguish a **public key**  $PK(data)$  and a **private (secret) key**  $SK(data)$ .

- Encrypt message from *Alice* to *Bob*:  $data = \underbrace{SK_{bob}(PK_{bob}(data))}_{\text{Action by Bob}}$  Sent by Alice
- Sign message for *Bob* by *Alice*:  
 $[data, \underbrace{data \stackrel{?}{=} PK_{alice}(SK_{alice}(data))}_{\text{Check by Bob}}] = [data, \underbrace{SK_{alice}(data)}_{\text{Sent by Alice}}]$

# Security mechanisms

## Secure hashing

In practice, we use **secure hash functions**:  $H(data)$  returns a **fixed-length string**.

- Any change from  $data$  to  $data^*$  will lead to a **completely different string**  $H(data^*)$ .
- Given a hash value, it is computationally impossible to find a  $data$  with  $h = H(data)$

# Security mechanisms

## Secure hashing

In practice, we use **secure hash functions**:  $H(data)$  returns a **fixed-length string**.

- Any change from  $data$  to  $data^*$  will lead to a **completely different string**  $H(data^*)$ .
- Given a hash value, it is computationally impossible to find a  $data$  with  $h = H(data)$

## Practical digital signatures

Sign message for *Bob* by *Alice*:

$$[data, \underbrace{H(data) \stackrel{?}{=} PK_{alice}(sgn)}}_{\text{Check by Bob}}] = \underbrace{[data, H, sgn = SK_{alice}(H(data))]}_{\text{Sent by Alice}}$$

# Scale in distributed systems

## Observation

Many developers of modern distributed systems easily use the adjective “scalable” without making clear **why** their system actually scales.

# Scale in distributed systems

## Observation

Many developers of modern distributed systems easily use the adjective “scalable” without making clear **why** their system actually scales.

## At least three components

- Number of users or processes (**size scalability**)
- Maximum distance between nodes (**geographical scalability**)
- Number of administrative domains (**administrative scalability**)

# Scale in distributed systems

## Observation

Many developers of modern distributed systems easily use the adjective “scalable” without making clear **why** their system actually scales.

## At least three components

- Number of users or processes (**size scalability**)
- Maximum distance between nodes (**geographical scalability**)
- Number of administrative domains (**administrative scalability**)

## Observation

Most systems account only, to a certain extent, for size scalability. Often a solution: multiple powerful servers operating independently in parallel. Today, the challenge still lies in geographical and administrative scalability.

# Advantages of Distributed Systems

- Improved performance and scalability
- Increased fault tolerance and Reliability
- Enhanced reliability
- Enhanced Data Sharing
- Geographic flexibility
- Cost-effectiveness



# Disadvantages of Distributed Systems

- Complexity
- Increased cost of set-up
- Security vulnerabilities
- Data inconsistency
- Network and resource overload
- Difficulty in debugging and troubleshooting

# Application Areas of Distributed Systems

Distributed Systems are Every Where These days!

- Peer to Peer Networks
- E-Commerce
- Cloud Computing
- The Internet and Web Wide Web
- Online Gaming
- Video Conferencing Platforms
- Aircraft and Industrial Control Systems

END